

Data

Table of contents

What is this dataset?	1
Why this dataset?	1
Source	1
Preparation pipeline	1
Download the cleaned data	2
Schema at a glance	2
A peek at the data	2

What is this dataset?

The dataset captures **hazardous cargo (HNS — Hazardous and Noxious Substances) handled at Turkish ports during 2022**. It was published by the **Directorate General for Maritime Affairs** (T.C. Ulaştırma ve Altyapı Bakanlığı – Denizcilik Genel Müdürlüğü) and shared with us in the form of a pivot-table workbook prepared by the directorate’s intern team.

Each leaf record describes a single cargo movement at the resolution of:

- the **regional port authority** (e.g., *Aliağa Bölge Liman Başkanlığı*),
- the **facility** that handled the cargo (e.g., *Egegaz Aliağa LNG Terminal*),
- the **receiver / final consignee company**,
- the **cargo amount in metric tons**.

Why this dataset?

We chose it for three reasons:

1. **It’s real, operational, and Turkish.** No simulated toy data — the patterns we surface are about the actual industrial geography of Türkiye’s coastlines.
2. **It supports policy questions.** The Maritime Affairs office has flagged firms handling more than 20,000 tons of hazardous cargo as candidates for upcoming tax/regulatory measures. The dataset lets us put numbers on that conversation.
3. **It’s multi-dimensional.** Company $\times$ facility $\times$ port authority gives us room for descriptive, geographic, and threshold-based analyses without leaving a single source.

Source

Directorate General for Maritime Affairs. *2022 Hazardous Cargo Handling Records*. Internal compilation, accessed via project supervisor.

The original Excel workbook is bundled with this site as `data/HNS_2022.xlsx`, alongside the cleaned versions described below.

Preparation pipeline

The raw workbook is a pivot-table export with three nested categorical levels (port authority $\rightarrow$ facility $\rightarrow$ company) and built-in subtotal rows (“Toplam ...”). To analyse it we need a flat, tidy table.

The full preprocessing script lives at `R/prepare_data.R`. In short, it:

1. Reads the workbook with `readxl::read_excel()` and assigns clean column names: `port_authority`, `facility`, `company`, `cargo_mt`.
2. **Drops every subtotal row** (Toplam ...) at all three hierarchy levels.
3. Drops rows where the company or tonnage is missing.
4. Renames the (boş) placeholder facility (used at Güllük) to **Unspecified facility**.
5. Builds a `company_clean` field — uppercased, ASCII-folded, punctuation-stripped — for grouping near-duplicate company names.
6. Extracts a `port_authority_short` field by stripping the boilerplate “Bölge Liman Başkanlığı” suffix.
7. Joins a **province + ISO-2 code** lookup so we can map each port authority onto a Turkish province for the choropleth.

Re-run the script any time the source workbook changes:

```
source("R/prepare_data.R")
```

It writes:

- `data/HNS_2022_clean.RData` — the canonical R object.
- `data/HNS_2022_clean.csv` — the same table in CSV (UTF-8) for inspection / sharing.

Download the cleaned data

Tip

The cleaned R dataset is available here: [HNS_2022_clean.RData](#) (also available as [HNS_2022_clean.csv](#)).

Loading it back is a one-liner:

```
load("data/HNS_2022_clean.RData") # creates a data frame called `hns2022`
```

Schema at a glance

Column	Type	Description
<code>port_authority</code>	character	Full port-authority name (e.g., <i>Aliğa Bölge Liman Başkanlığı</i>)
<code>port_authority_short</code>	character	Short form (e.g., <i>Aliğa</i>), used for joins
<code>facility</code>	character	Facility / terminal that handled the cargo
<code>company</code>	character	Receiver / consignee, as written in the source
<code>company_clean</code>	character	Normalised company name for grouping near-duplicates
<code>cargo_mt</code>	numeric	Cargo amount, in metric tons
<code>province</code>	character	Host province of the port authority
<code>iso2</code>	character	ISO 3166-2 code for the province (e.g., TR-35)

A peek at the data

```
hns2022 |>
  dplyr::select(port_authority_short, facility, company, cargo_mt, province) |>
  head(10) |>
  kable(
    col.names = c("Port authority", "Facility", "Company", "Cargo (m. ton)", "Province"),
    format.args = list(big.mark = ",")
  )
```

```
) |>
kable_styling(bootstrap_options = c("striped", "hover"), full_width = FALSE)
```

Table 2: First 10 rows of the cleaned dataset.

Port authority	Facility	Company	Cargo (m. ton)	Province
Aliğa	TÜPRAŞ	İzmir Aliğa Rafinerisi	15,088,213.00	İzmir
Aliğa	STAR ALIAGA	STAR RAFİNERİ A.Ş	12,114,504.29	İzmir
Aliğa	STAR ALIAGA	SOCAR	1,173,424.00	İzmir
Aliğa	STAR ALIAGA	PETKİM	222,422.88	İzmir
Aliğa	STAR ALIAGA	POAŞ	93,866.95	İzmir
Aliğa	STAR ALIAGA	SHELL/ISTANBUL JET/SOCAR	67,807.00	İzmir
Aliğa	STAR ALIAGA	STAR	65,734.57	İzmir
Aliğa	STAR ALIAGA	OMV	49,875.76	İzmir
Aliğa	STAR ALIAGA	SOCAR,SHELL,THY OPET	49,211.41	İzmir
Aliğa	STAR ALIAGA	THY OPET	32,855.83	İzmir